

Work progress

Development of:

- ▶ A generic specification template for trace validation (called trace spec.)
- ▶ A library (called instrumentation) that enable to log events and changes happening on variables
- ▶ A “method” that aims to log implementation properly

3 implementations:

- ▶ Two phase protocol (distributed)
- ▶ Key value store
- ▶ Raft (distributed)

Raft example - spec

Here an example of a base specification (Raft):

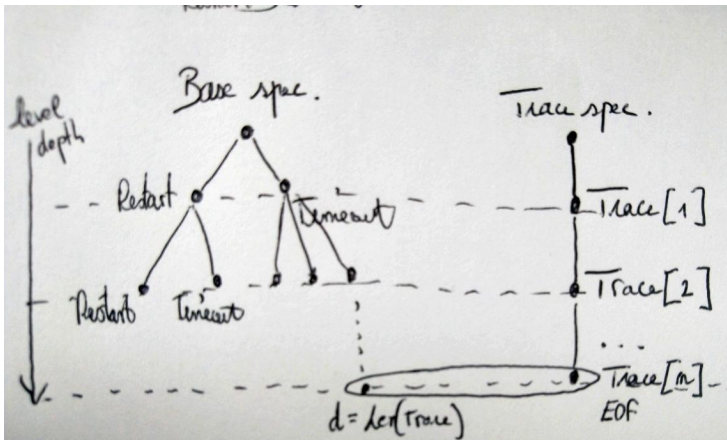
```
\* Defines how the variables may transition.  
Next == /\ \/\E i \in Server : Restart(i)  
        \/\E i \in Server : Timeout(i)  
        \/\E i \in Server : BecomeLeader(i)  
  
...  
...  
  
Spec == Init /\ [] [Next]_vars
```

Raft example - trace

- ▶ A trace can be seen as a behavior of a system
- ▶ A trace is a sequence of events (atomic TLA+ action)
- ▶ Each event is compound by one or many variable updates
- ▶ Below an extract of a trace of raft consensus algorithm

```
{
  "clock": 1,
  "state": [
    {"op": "Replace", "path": ["node2"],
     "args": ["Candidate"]}],
  "desc": "Timeout"
}
...
{
  "clock": 26,
  "state": [{"op": "Replace", "path": ["node1"],
             "args": ["Leader"]}],
  "desc": "BecomeLeader"
}
```

Trace specification - how validate a trace ?



Trace specification - how validate a trace ?

- ▶ At least one path of state space graph must lead to the complete reading of trace
- ▶ Use a POSTCONDITION (hyperproperty)
- ▶ Allow TLC to have non-deterministic behavior

```
TraceAccepted ==
```

```
  (* Diameter equal to trace length => *)
```

```
  (* Trace file has been read completely at least one time *)
```

```
  LET d == TLCGet("stats").diameter IN
```

```
  IF d - 1 = Len(Trace) THEN TRUE
```

```
  ELSE Print(<<"Failed matching the trace to (a prefix of "
```

```
    "TLA+ debugger breakpoint hit count " \
```

```
POSTCONDITION
```

```
  TraceAccepted
```

Trace specification - how it work ? - spec refinement

- We have to write a trace spec that is a refinement of a base spec (here Raft)

```
(* Temporal formula for trace spec *)
```

```
TraceSpec == TraceInit /\ [] [TraceNext]_<<1, vars>>
```

```
(* Instanciate raft *)
```

```
BASE == INSTANCE raft
```

```
BaseSpec == BASE!Init /\ [] [BASE!Next \/ ComposedNext]_BASE
```

```
SPECIFICATION
```

```
    TraceSpec
```

```
PROPERTIES
```

```
    (* Refine raft *)
```

```
    BaseSpec
```

Trace specification - read trace events

- ▶ Read trace line after line (each line is an event)
- ▶ Apply all operations, on all variables found in each events

```
logline == Trace[1]
```

```
ReadNext ==
```

```
  (* depth: line number *)  
  /\ l' = l + 1  
  (* Apply all variable updates *)  
  /\ MapVariables(logline)  
  (* Advance base spec *)  
  /\ BaseSpec::Next
```

Trace specification - variables update and mapping

- ▶ TLC apply all operations to all variables precised in current event

```
MapVariables(logline) ==  
  /\  
    IF "state" \in DOMAIN logline  
    THEN state' = ExceptAtPaths(state, "state", logline)  
    ELSE TRUE  
  /\  
    IF "currentTerm" ...
```

Note: If a variable changes isn't logged, TraceSpec just let TLC search for all possible values of this variable according to base spec (see TRUE).

Trace specification - variables update and mapping

- ▶ Variable updates was made by applying 1 or more operators on it
- ▶ Operators are generic and defined in trace spec, for example:

```
Replace(cur, val) == val
```

```
AddElement(cur, val) == cur \cup {val}
```

```
AddElements(cur, vals) == cur \cup ToSet(vals)
```

```
RemoveElement(cur, val) == cur \ {val}
```

```
Clear(cur, val) == {}
```

```
...
```

Trace specification - variables update and mapping

- ▶ following event:

```
{  
  "clock": 1,  
  "state": [  
    {"op": "Replace", "path": ["node2"],  
     "args": ["Candidate"]}],  
  "desc": "Timeout"  
}
```

- ▶ should map variable state as following:

```
state' = [state EXCEPT !["node2"] = "Candidate"]
```

Trace specification - variables update and mapping

- ▶ A variable can be updated partially at a given path

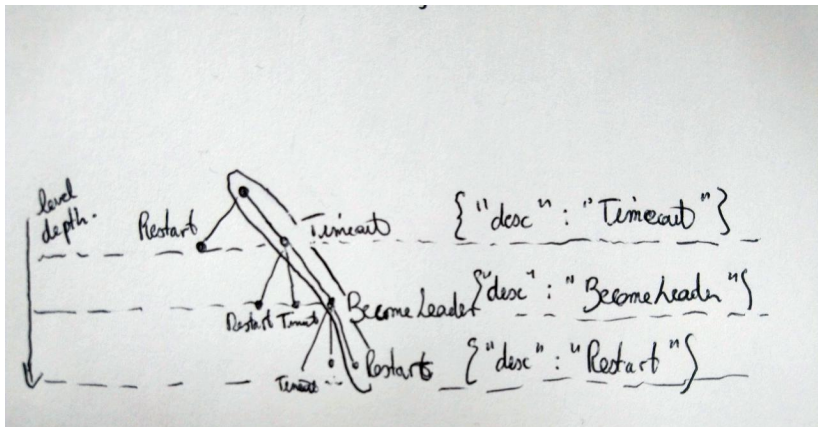
```
{"matchIndex": [{  
  "op": "Replace",  
  "path": ["node3", "node2"],  
  "args": [7]}]}
```

- ▶ This update will be automatically translated to:

```
matchIndex' = [matchIndex EXCEPT !["node3"] ["node2"] = 7]
```

Trace specification - optimisation

- ▶ State space can be largely reduced if we precise the name of the next action expected. Action name lead TLC and select directly the expected action. Therefore reducing non-deterministic behaviors.
- ▶ Specify action name when logging is recommended but not mandatory



Trace specification - optimisation

- ▶ For each action contained in base spec we write a corresponding predicate
- ▶ Predicate enable TLC to select next expected action when `IsEvent` is `TRUE`

```
IsEvent(e) ==  
  /\ IF "desc" \in DOMAIN logline  
    THEN logline.desc = e ELSE TRUE
```

```
IsRestart ==  
  /\ IsEvent("Restart")  
  /\ \E i \in Server : Restart(i)
```

```
IsTimeout ==  
  /\ IsEvent("Timeout")  
  /\ \E i \in Server : Timeout(i)
```

...

Trace specification - optimisation

- ▶ next action of trace spec is just the disjunction of all predicates

```
TraceNext ==  
  \/ IsRestart  
  \/ IsTimeout  
  ...
```

Instrumentation - purpose

- ▶ Aims to generate a trace by logging some events
- ▶ Aims to log event and variable changes

Trace example:

```
{  
  "clock": 1,  
  "state": [{"op": "Replace", "path": ["node2"], "args":  
    "commitIndex": [{"op": "Replace", "path": ["node2"], "a  
  "desc": "Restart"  
}  
...
```

Instrumentation - How to log

1. We have to log events: log all commits is necessary because TLC cannot fill holes in events
2. We have to log variable changes: log of all variables isn't necessary, but more variables we log, more the statespace reduce, and more we are confident in the implementation

Instrumentation - log event

Example of log “Timeout” event in Raft:

```
public void timeout() {  
    assert state == NodeState.Follower;  
    ...  
    spec.commitChanges("Timeout");  
}
```

Instrumentation - log variables

The idea is to log variable updates like you manipulate directly the specification's variables.

Declare spec variable example:

```
this.spec = new TraceInstrumentation(nodeInfo.name() + ".no
// Binding to variable state at path nodeName (state[nodeN
this.specState = spec.getVariable("state")
                        .getField(nodeInfo.name());
this.specVotesGranted = spec.getVariable("votesGranted")
                        .getField(nodeInfo.name());
```

Instrumentation - log variables

Log variable changes example:

```
private void setState(NodeState state) {  
    this.state = state;  
    // this.spec.notify(specState, SET, state.toString());  
    specState.set(state.toString());  
}  
  
...  
if (m.isGranted()) {  
    // Add node that granted a vote to me  
    candidateState.getGranted().add(m.getFrom());  
    specVotesGranted.add(m.getFrom());  
}
```

Instrumentation - clocks

We can use two way to sync clock between distributed processes:

- ▶ Lamport clock, we send clock in the message and we call explicitly sync method on logging framework
- ▶ Shared clock, if all the system is executed on the same physical machine, all process can share a clock in a memory mapped file: `SharedClock.get(clockName);`

Execution pipeline

In all our tests we make a script execution pipeline that do the following:

- ▶ Execute implementation (which create a trace file by logging events and variable updates)
- ▶ Merge trace files that was produced by different processes
- ▶ Execute TLC on the trace spec for a given trace file in order to make validation