

# Work progress

Development of:

- ▶ A generic specification template for trace validation (*trace spec*)
- ▶ A library (*instrumentation*) for logging events and variable updates
- ▶ A “method” based on the above for validating traces of implementations

Applied to three case studies:

- ▶ Two-phase protocol (distributed)
- ▶ Key-value store
- ▶ Raft (distributed)

## Raft example - spec

The base specification (Raft):

```
\* Defines the system transitions
Next == /\ \/ \E i \in Server : Restart(i)
          \/ \E i \in Server : Timeout(i)
          \/ \E i \in Server : BecomeLeader(i)
...
...

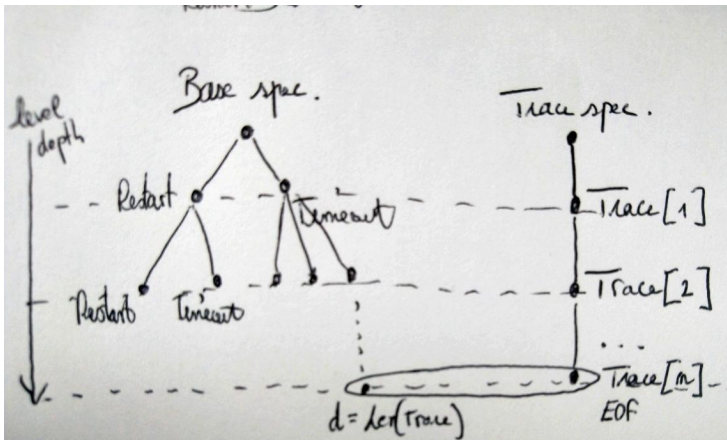
Spec == Init /\ [] [Next]_vars
```

## Raft example - trace

- ▶ A trace records a behavior of a system
- ▶ A trace is a sequence of events (corresponding to TLA+ actions)
- ▶ Each event may contain several variable updates
- ▶ Extract of a trace of Raft:

```
{
  "clock": 1,
  "state": [ {"op": "Replace",
              "path": ["node2"],
              "args": ["Candidate"]} ],
  "desc": "Timeout"
}
...
{
  "clock": 26,
  "state": [ {"op": "Replace",
              "path": ["node1"],
              "args": ["Leader"]} ],
  "desc": "BecomeLeader"
}
```

## Trace specification - how do we validate a trace ?



## Trace specification - how do we validate a trace ?

- ▶ The trace must correspond to at least one path through the state space graph
- ▶ Expressed for TLC as a POSTCONDITION
- ▶ Exploit non-determinism of TLA+ specifications

```
TraceAccepted ==
```

```
  (* Diameter equal to trace length => *)
```

```
  (* Trace file has been read completely at least one time *)
```

```
  LET d == TLCGet("stats").diameter IN
```

```
  IF d - 1 = Len(Trace) THEN TRUE
```

```
  ELSE Print(<<"Failed matching the trace to (a prefix of  
    "TLA+ debugger breakpoint hit count " \
```

```
POSTCONDITION
```

```
  TraceAccepted
```

## Ensure that the trace specification refines the base spec

```
(* Temporal formula for trace spec *)
TraceSpec == TraceInit /\ [] [TraceNext]_<<1, vars>>

(* Instantiate raft *)
BASE == INSTANCE raft
BaseSpec == BASE!Init /\ [] [BASE!Next \/ ComposedNext]_BASE!vars

SPECIFICATION
    TraceSpec
PROPERTIES
    (* Refine raft *)
    BaseSpec
```

## Trace specification - read trace events

- ▶ Read trace one line at a time (each line is an event)
- ▶ Update variables according to information provided by events

```
logline == Trace[l]
```

```
ReadNext ==
```

```
  (* depth: line number *)  
  /\ l' = l + 1  
  (* Apply all variable updates *)  
  /\ MapVariables(logline)  
  (* Advance base spec *)  
  /\ BASE!Next
```

## Trace specification - updating variables

```
MapVariables(logline) ==  
  /\  
    IF "state" \in DOMAIN logline  
    THEN state' = ExceptAtPaths(state, "state",  
                                logline.state)  
    ELSE TRUE  
  /\  
    IF "currentTerm" ...
```

If a variable change isn't logged, TraceSpec just lets TLC search for all possible values of this variable according to base spec.



# Trace specification - updating variables

- Generic operators for updating variables

`Replace(cur, val) == val`

`AddElement(cur, val) == cur \cup {val}`

`AddElements(cur, vals) == cur \cup vals`

`RemoveElement(cur, val) == cur \ {val}`

`Clear(cur, val) == {}`

`...`

# Trace specification - updating variables

The event

```
{  
  "clock": 1,  
  "state": [ {"op": "Replace",  
              "path": ["node2"],  
              "args": ["Candidate"]} ],  
  "desc": "Timeout"  
}
```

should map the variable state as follows:

```
state' = [state EXCEPT !["node2"] = "Candidate"]
```

# Trace specification - updating variables

- ▶ A variable can be updated partially at a given path

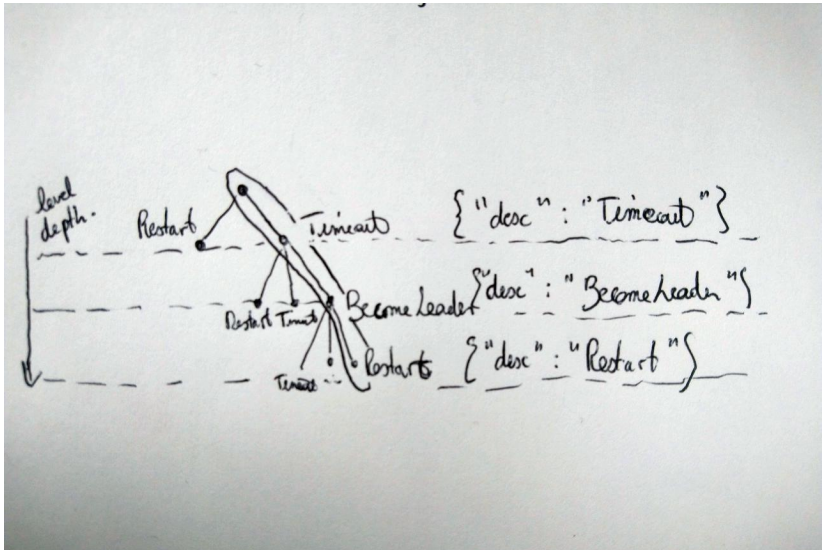
```
{"matchIndex": [ {"op": "Replace",  
                  "path": ["node3", "node2"],  
                  "args": [7]}]}
```

- ▶ This update will be automatically translated to:

```
matchIndex' = [matchIndex EXCEPT !["node3"] ["node2"] = 7]
```

# Trace specification - optimization

- ▶ In order to reduce the state space, the trace may indicate the name of the action to be applied.
- ▶ Specifying action name when logging is not mandatory.



## Trace specification - optimization

- ▶ For each action contained in base spec we write a corresponding predicate
- ▶ Predicate enable TLC to select next expected action when IsEvent is TRUE

```
IsEvent(e) ==  
  /\ IF "desc" \in DOMAIN logline  
    THEN logline.desc = e ELSE TRUE
```

```
IsRestart ==  
  /\ IsEvent("Restart")  
  /\ \E i \in Server : Restart(i)
```

```
IsTimeout ==  
  /\ IsEvent("Timeout")  
  /\ \E i \in Server : Timeout(i)
```

...

# Trace specification - optimization

- ▶ next action of trace spec is just the disjunction of all predicates

```
TraceNext ==  
  \/ IsRestart  
  \/ IsTimeout  
  ...
```

# Instrumentation - purpose

- ▶ Generate a trace by logging some events
- ▶ Log event and variable changes

Trace example:

```
{  
  "clock": 1,  
  "state": [{"op": "Replace", "path": ["node2"], "args":  
    "commitIndex": [{"op": "Replace", "path": ["node2"], "a  
    "desc": "Restart"  
}  
...
```

# Instrumentation - How to log

1. We have to log all events that correspond to actions of the base spec: TLC will not fill “holes”.
2. Logging all variable updates is not necessary, but the more variables we log, the smaller is the state space explored by TLC, and the more confident we are in the implementation



# Instrumentation - log events

Example of log “Timeout” event in Raft:

```
public void timeout() {  
    assert state == NodeState.Follower;  
    ...  
    spec.commitChanges("Timeout");  
}
```

# Instrumentation - logging variables

The idea is to log variable updates whenever a variable corresponding to a specification variable is modified.

Declare spec variable:

```
this.spec = new TraceInstrumentation(nodeInfo.name()  
                                     + ".ndjson", clock);  
// Binding to variable state at path nodeName (state[nodeName ])  
this.specState = spec.getVariable("state")  
                    .getField(nodeInfo.name());  
this.specVotesGranted = spec.getVariable("votesGranted")  
                           .getField(nodeInfo.name());
```

# Instrumentation - logging variables

Log variable changes:

```
private void setState(NodeState state) {  
    this.state = state;  
    // this.spec.notify(specState, SET, state.toString());  
    specState.set(state.toString());  
}  
  
...  
if (m.isGranted()) {  
    // Add node that granted a vote to me  
    candidateState.getGranted().add(m.getFrom());  
    specVotesGranted.add(m.getFrom());  
}
```

# Instrumentation - clocks

Two ways of synchronizing clock between distributed processes are supported:

- ▶ Lamport clock: clocks are sent in messages and we explicitly call the sync method of logging framework
- ▶ Shared clock, if all processes are executed on the same physical machine, they can share a clock in a memory mapped file:  
`SharedClock.get(clockName);`

# Execution pipeline

Tests are run as a script execution pipeline:

- ▶ Execute implementation (this creates a trace file by logging events and variable updates)
- ▶ Merge trace files that were produced by different processes
- ▶ Execute TLC on the trace spec for a given trace file

## Results: bugs found

- ▶ KeyValueStore: forgotten conditions / guards (3 cases)
- ▶ Raft: strict instead of non-strict inequalities
- ▶ Instrumentation: forgotten thread synchronisation
- ▶ Bugs can be identified very quickly:
  - ▶ Use of desc field gives information about the action that failed
  - ▶ Retrieve line number where validation fails and use TLA+ debugger

## Results: benefits and limits

- ▶ Find bugs in new implementations: events adhere to the specification
- ▶ Avoid regressions when implementation changes
- ▶ Need to know the specification
  - ▶ Especially all the actions (to be able to log all events)
  - ▶ The structure of variables (to be able to update them partially)
  - ▶ The part of the system that is distributed